

02_repr_tradicionales

October 27, 2025

```
[1]: # Generar representaciones tradicionales:
# - TF-IDF de palabras.
# - TF-IDF de n-gramas de caracteres.
# Usaremos solo TRAIN para ajustar y transformaremos TRAIN y VALIDATION.
# Guardamos matrices .npz, vectorizers .pkl y un índice con doc_id/sent_id.
```

Imports y config

```
[2]: from pathlib import Path
import pandas as pd
import numpy as np
import pickle, json
from scipy import sparse
from sklearn.feature_extraction.text import TfidfVectorizer

pd.set_option("display.max_colwidth", 120)
SEED = 42
```

Rutas

```
[3]: def find_root():
    p = Path.cwd()
    for cand in [p, *p.parents]:
        if (cand / "data" / "processed").exists():
            return cand
    raise FileNotFoundError("No encuentro data/processed.")

ROOT = find_root()
PROC = ROOT / "data" / "processed"
FEAT = ROOT / "features" / "tfidf"
FEAT.mkdir(parents=True, exist_ok=True)

NIVELES = ["easy", "medium", "hard"]
SPLITS = ["train", "validation"]
```

Carga de processed → DataFrame

```
[4]: def cargar_processed():
    filas = []
    for level in NIVELES:
```

```

    for split in SPLITS:
        p = PROC / level / split / "sentences.jsonl"
        if not p.exists():
            continue
        df = pd.read_json(p, lines=True)
        df["level"] = level
        df["split"] = split
        filas.append(df[["doc_id", "sent_id", "level", "split", "text_norm"]])
    return pd.concat(filas, ignore_index=True)

df = cargar_processed().sort_values(["level", "split", "doc_id", "sent_id"]).
    ↪reset_index(drop=True)
print(df.shape, "filas")
df.head(3)

```

(208160, 5) filas

```

[4]:      doc_id  sent_id level  split \
0  problem-1         0  easy  train
1  problem-1         1  easy  train
2  problem-1         2  easy  train

```

```

                                text_norm
0  there s also incidents of testosterone insensitive males that have either
ambiguous genitals or are phenotypically f...
1                                female is the human default body plan so a number of
conditions exist that cause female looking males
2                                what about people born appearing num female
complete with the bits but are genetically males

```

TF-IDF de palabras

```

[5]: # Ajuste en TRAIN
df_train = df[df["split"] == "train"]
corpus_train = df_train["text_norm"].astype(str).tolist()

vec_word = TfidfVectorizer(
    lowercase=False,          # ya normalizado
    analyzer="word",
    ngram_range=(1,2),       # unigrams + bigrams
    min_df=5,                # filtra ruido
    max_features=100_000
)
Xtr_w = vec_word.fit_transform(corpus_train)

# Transform VALIDATION
df_val = df[df["split"] == "validation"]
Xva_w = vec_word.transform(df_val["text_norm"].astype(str).tolist())

```

```

print("WORD TF-IDF")
print("train:", Xtr_w.shape, "nnz:", Xtr_w.nnz)
print("val  :", Xva_w.shape, "nnz:", Xva_w.nnz)

```

WORD TF-IDF

train: (171602, 85932) nnz: 4449125

val : (36558, 85932) nnz: 937263

TF-IDF de n-gramas de caracteres

```

[6]: vec_char = TfidfVectorizer(
        analyzer="char",
        ngram_range=(3,5),
        min_df=5,
        max_features=200_000
    )
Xtr_c = vec_char.fit_transform(corpus_train)
Xva_c = vec_char.transform(df_val["text_norm"].astype(str).tolist())

print("CHAR TF-IDF")
print("train:", Xtr_c.shape, "nnz:", Xtr_c.nnz)
print("val  :", Xva_c.shape, "nnz:", Xva_c.nnz)

```

CHAR TF-IDF

train: (171602, 134419) nnz: 45449143

val : (36558, 134419) nnz: 9754529

Guardado de matrices y metadatos

```

[7]: # Índices para mapear filas + (level, split, doc_id, sent_id)
idx_train = df_train[["level", "split", "doc_id", "sent_id"]].
    ↪reset_index(drop=True)
idx_val    = df_val[["level", "split", "doc_id", "sent_id"]].reset_index(drop=True)

# Carpetas
( FEAT / "word" ).mkdir(parents=True, exist_ok=True)
( FEAT / "char" ).mkdir(parents=True, exist_ok=True)

# WORD
sparse.save_npz(FEAT / "word" / "X_train_word.npz", Xtr_w)
sparse.save_npz(FEAT / "word" / "X_val_word.npz",   Xva_w)
pickle.dump(vec_word, open(FEAT / "word" / "vectorizer_word.pkl", "wb"))
idx_train.to_csv(FEAT / "word" / "index_train.csv", index=False)
idx_val.to_csv(  FEAT / "word" / "index_val.csv",   index=False)

# CHAR
sparse.save_npz(FEAT / "char" / "X_train_char.npz", Xtr_c)
sparse.save_npz(FEAT / "char" / "X_val_char.npz",   Xva_c)

```

```

pickle.dump(vec_char, open(FEAT / "char" / "vectorizer_char.pkl", "wb"))
idx_train.to_csv(FEAT / "char" / "index_train.csv", index=False)
idx_val.to_csv(FEAT / "char" / "index_val.csv", index=False)

# Informe JSON mínimo
reporte = {
    "word": {
        "shape_train": Xtr_w.shape, "nnz_train": int(Xtr_w.nnz),
        "shape_val": Xva_w.shape, "nnz_val": int(Xva_w.nnz),
        "vocab_size": len(vec_word.vocabulary_)
    },
    "char": {
        "shape_train": Xtr_c.shape, "nnz_train": int(Xtr_c.nnz),
        "shape_val": Xva_c.shape, "nnz_val": int(Xva_c.nnz),
        "vocab_size": len(vec_char.vocabulary_)
    }
}
(Path(FEAT) / "tfidf_resumen.json").write_text(json.dumps(reporte, indent=2),
encoding="utf-8")
print("Guardado en features/tfidf/")

```

Guardado en features/tfidf/

Tests Finales

```

[8]: # Terminos más pesados por TF-IDF medio en TRAIN (palabras)
col_means = np.asarray(Xtr_w.mean(axis=0)).ravel()
top_idx = col_means.argsort()[-10:][:-1]
inv_vocab = {j:i for i,j in vec_word.vocabulary_.items()}
top_terms = [inv_vocab[i] for i in top_idx]
pd.DataFrame({"term": top_terms, "tfidf_mean": col_means[top_idx].round(6)})

```

```

[8]:
   term  tfidf_mean
0   the    0.034300
1    to    0.026840
2   and    0.022177
3    of    0.020698
4    in    0.018026
5    it    0.017368
6    is    0.017342
7   that  0.016891
8   they  0.013748
9    you  0.013222

```

1 Informe breve — 02_representaciones_tradicionales.ipynb

1.1 Objetivo

Crear representaciones TF-IDF sobre frases normalizadas en `processed` para alimentar baselines y modelos posteriores.

1.2 Configuración

- Ajuste en **train**. Transformación en **train** y **validation**.
- Dos variantes:
 - Palabras: `ngram_range=(1,2)`, `min_df=5`, `max_features=100k`, `lowercase=False`.
 - Caracteres: `ngram_range=(3,5)`, `min_df=5`, `max_features=200k`.

1.3 Salidas

- `features/tfidf/word/`
 - `X_train_word.npz`, `X_val_word.npz`
 - `vectorizer_word.pkl`
 - `index_train.csv`, `index_val.csv`
- `features/tfidf/char/`
 - `X_train_char.npz`, `X_val_char.npz`
 - `vectorizer_char.pkl`
 - `index_train.csv`, `index_val.csv`
- Resumen: `features/tfidf/tfidf_resumen.json` con shapes, nnz y vocabulario.

1.4 Uso previsto

- Comparar baselines de cambio de autor calculando distancias coseno entre frases o ventanas.
- Servir de entrada a clasificadores ligeros en E3/E4.

1.5 Notas

- El texto ya viene normalizado desde 01_preprocesamiento.
- Los índices CSV permiten mapear cada fila a (`level`, `split`, `doc_id`, `sent_id`).
- Los hiperparámetros están pensados para un primer pase. Ajustables según memoria y rendimiento.